

Fine-tuning

Tirtharaj Dash

Department of CS & IS
BITS Pilani, Goa Campus, India
[Homepage](#)

February 2026

Outline

Recap: sampling from the LLM distribution

What is θ ? The GPT block

Fine-tuning: adapting θ to a new distribution

Limitations: distribution shift and catastrophic forgetting

Parameter-efficient fine-tuning (PEFT) and LoRA

API-based fine-tuning

The LLM as a probability distribution

An LLM defines a joint distribution over token sequences:

$$P(\mathbf{x}; \theta) = P(x_1, \dots, x_T; \theta) = \prod_{t=1}^T P(x_t | x_1, \dots, x_{t-1}; \theta)$$

Generation = sampling from this distribution.

Given a prompt (x_1, \dots, x_s) , we draw the next token:

$$x_{s+1} \sim P(\cdot | x_1, \dots, x_s; \theta)$$

and repeat until an end-of-sequence token is produced.


The LLM as a probability distribution

- ▶ *Greedy decoding*: always pick the most probable token
 $x_t = \arg \max P(\cdot | x_{<t}; \theta)$
- ▶ *Temperature sampling*: soften the distribution before sampling, controlling randomness
- ▶ *Top-p / top-k sampling*: restrict sampling to the most probable tokens

The model knows *nothing* beyond what is encoded in θ .

Inside a GPT block: what θ actually is

A GPT-style LLM stacks L identical transformer decoder blocks.
Each block contains:



θ is the collection of all weight matrices across all L blocks:

$$\theta = \{W_Q^{(\ell)}, W_K^{(\ell)}, W_V^{(\ell)}, W_O^{(\ell)}, W_1^{(\ell)}, W_2^{(\ell)}, \gamma^{(\ell)}, \beta^{(\ell)}\}_{\ell=1}^L$$

A 7B model has ~ 7 billion such scalar values.

Pre-training: learning θ from data

Objective: minimise the negative log-likelihood over a massive corpus \mathcal{D}_{pre} :

$$\theta_0 = \arg \min_{\theta} -\mathbb{E}_{x \sim \mathcal{D}_{\text{pre}}} \left[\sum_{t=1}^T \log P(x_t \mid x_{<t}; \theta) \right]$$

- ▶ \mathcal{D}_{pre} : trillions of tokens from the web, books, code, etc.
 - ▶ After training, θ_0 encodes broad linguistic knowledge, world facts, and reasoning patterns
 - ▶ The resulting model is a *generalist*: it can continue any text, but does not reliably follow instructions or specialise in any domain
- ⇒ θ_0 is the starting point. Fine-tuning adjusts it.

Fine-tuning as distribution shift

Key idea: the pre-training corpus \mathcal{D}_{pre} defines one data distribution. A downstream task defines a different, narrower distribution \mathcal{D}_{ft} .

Fine-tuning continues gradient descent from θ_0 on \mathcal{D}_{ft} :

$$\theta^* = \arg \min_{\theta} -\frac{1}{N} \sum_{i=1}^N \sum_t \log P(y_t^{(i)} | x^{(i)}, y_{<t}^{(i)}; \theta)$$

where $(x^{(i)}, y^{(i)})$ are (input, target) pairs and the loss is computed *only on the output tokens y* .

Fine-tuning as distribution shift

Pre-training distribution

General web text, books, code.
 $|\mathcal{D}_{\text{pre}}| \sim 10^{12}$ tokens.

Fine-tuning distribution

Curated (instruction, response) pairs.
 $|\mathcal{D}_{\text{ft}}| \sim 10^3\text{--}10^5$ examples.

The model retains representations from θ_0 while its output distribution shifts to match \mathcal{D}_{ft} .

Two failure modes of fine-tuning

Distribution shift at test time

If the fine-tuning distribution \mathcal{D}_{ft} differs from what the model will see at deployment, performance degrades.

$$\mathcal{D}_{\text{ft}} \neq \mathcal{D}_{\text{test}} \Rightarrow \text{generalisation gap}$$

Example: fine-tuned on clinical notes from one hospital system; deployed at another with different terminology and documentation style.

Two failure modes of fine-tuning



Catastrophic forgetting

Gradient updates that minimise loss on \mathcal{D}_{ft} may increase loss on \mathcal{D}_{pre} . The model *forgets* general capabilities.

$$\theta^* \approx \arg \min \mathcal{L}_{\text{ft}} \not\Rightarrow \theta^* \approx \arg \min \mathcal{L}_{\text{pre}}$$

The two objectives can conflict, especially with a large learning rate or many fine-tuning steps.

Kirkpatrick et al., *Overcoming Catastrophic Forgetting*, PNAS 2017

Why not update all of θ ?

For a model with P parameters, full fine-tuning requires storing in GPU memory:

| What | Memory (fp32) | 7B model |
|--|-------------------------------|---------------------------------|
| Model weights θ | $4P$ bytes | 28 GB |
| Gradients $\nabla_{\theta}\mathcal{L}$ | $4P$ bytes | 28 GB |
| Adam optimiser states | $8P$ bytes | 56 GB |
| Total | $16P$ bytes | > 112 GB |

This exceeds a single A100 GPU (80 GB). Even in fp16 it is borderline.

Why not update all of θ ?

PEFT idea: freeze θ_0 ; introduce a small set of trainable parameters ϕ with $|\phi| \ll |\theta|$:

$$\theta^* = \underbrace{\theta_0}_{\text{frozen}} + \underbrace{\Delta\theta(\phi)}_{\text{trained}}$$

Train only ϕ . Gradient and optimiser memory scales with $|\phi|$, not $|\theta|$.

LoRA: the low-rank hypothesis

Empirical observation (Aghajanyan et al., 2021): the weight updates $\Delta\theta$ that arise during fine-tuning have low *intrinsic rank* — the update lives in a much smaller subspace than the full weight space.

LoRA (Hu et al., 2022) exploits this. For each weight matrix $W_0 \in \mathbb{R}^{d \times k}$, instead of learning a full ΔW , learn a rank- r factorisation:

$$W = W_0 + \Delta W = W_0 + BA$$

where $B \in \mathbb{R}^{d \times r}$, $A \in \mathbb{R}^{r \times k}$, $r \ll \min(d, k)$.

LoRA: the low-rank hypothesis

- ▶ W_0 is *frozen*; only A and B are trained
- ▶ A initialised with random Gaussian; B initialised to **zero** $\Rightarrow \Delta W = 0$ at the start (no disruption to pre-trained behaviour)
- ▶ Forward pass: $h = W_0x + \frac{\alpha}{r} BAx$ (α: scaling hyperparameter)

Hu et al., *LoRA: Low-Rank Adaptation of Large Language Models*, ICLR 2022

LoRA: the low-rank hypothesis

LoRA: parameter count:

For a single weight matrix $W_0 \in \mathbb{R}^{d \times k}$:

| | Full fine-tuning | LoRA (rank r) |
|--------------------------------|------------------|-------------------|
| Parameters | $d \times k$ | $r(d + k)$ |
| Example ($d=k=4096$, $r=8$) | 16.7M | 65K |
| Reduction | — | $\sim 256 \times$ |

LoRA: the low-rank hypothesis

After training: **merge** the adapter back into the base weight —


$$W^* = W_0 + BA$$

No inference overhead; the merged model is identical in structure to the original.

Typical choices: $r \in \{4, 8, 16, 64\}$; larger r = more capacity but more parameters. For most tasks $r = 8$ suffices.

LoRA applied to each GPT block

LoRA adapters are inserted into the *linear projection matrices* of each transformer block. Typically applied to the attention projections:



- ▶ All blue (attention) and red (FFN) weights are **frozen**
- ▶ Green LoRA adapters (B, A) are **trained**
- ▶ This is repeated independently for each of the L blocks

Fine-tuning without touching the model

Not everyone can run fine-tuning locally. Cloud providers expose fine-tuning as an API:

| Provider | Models | Method (reported) |
|--------------|-------------------------|------------------------------|
| OpenAI | GPT-4o mini, GPT-3.5 | SFT (likely LoRA internally) |
| Google | Gemini 1.5 Flash | SFT + RLHF |
| Anthropic | Claude (limited access) | Constitutional AI + RLHF |
| Together AI | LLaMA, Mistral, etc. | LoRA / QLoRA |
| Hugging Face | Any open model | Full FT / LoRA via AutoTrain |

Workflow: upload \mathcal{D}_{ft} in JSONL format (one {prompt, completion} per line) → submit job → receive a fine-tuned model endpoint.

Fine-tuning without touching the model

Trade-offs vs. local fine-tuning

- ▶ No GPU needed; scales to large models; simple API
- ▶ Data leaves your machine; no control over PEFT rank or hyperparameters; cost per token at inference; model not portable

Summary

- ▶ An LLM defines $P(\mathbf{x}; \theta)$; generation is sampling from this distribution; θ is all weight matrices in all L blocks
- ▶ Pre-training learns θ_0 from massive general data; fine-tuning shifts it to θ^* on small task-specific data
- ▶ Fine-tuning loss: cross-entropy on *output tokens only*
- ▶ **Distribution shift**: mismatch between \mathcal{D}_{ft} and deployment distribution causes generalisation failure
- ▶ **Catastrophic forgetting**: \mathcal{L}_{ft} and \mathcal{L}_{pre} objectives can conflict
- ▶ **LoRA**: approximate $\Delta W \approx BA$ with rank r ; freeze W_0 , train only B and A ; merge at inference
- ▶ LoRA is applied independently to each GPT block's projection matrices
- ▶ **API fine-tuning**: upload data, get a fine-tuned endpoint — convenient but at the cost of control and data privacy