

LLM with Retrieval-Augmented Generation

Tirtharaj Dash

Department of CS & IS
BITS Pilani, Goa Campus, India
[Homepage](#)

February 2026

Outline

Recap: LLMs and their limitations

Retrieval-Augmented Generation (RAG)

The retrieval function

Generation with retrieved context

Evaluation of RAG systems

Summary: RAG interface

Demo: Jupyter

LLMs: what we know so far

An autoregressive LLM factorises the joint distribution over a token sequence $\mathbf{x} = (x_1, x_2, \dots, x_T)$ as:

$$P(\mathbf{x}; \theta) = \prod_{t=1}^T P(x_t | x_1, \dots, x_{t-1}; \theta)$$

- ▶ The model parameters θ compress all knowledge seen during training
- ▶ At inference: given a prompt q , the model generates

$$y^* = \arg \max_y P(y | q; \theta)$$

- ▶ Generation is purely *parametric*: no external lookup, no memory

Three fundamental limitations

Stale knowledge

θ is fixed after training. Any fact that changes or emerges after the training cutoff is inaccessible to the model.

Hallucination

$P(y | q; \theta)$ is a *probability distribution* over tokens, not a lookup table. The model generates fluent text even when it has no reliable basis for the claim.

No attribution

There is no mechanism to trace *which* training document caused a particular output making verification impossible.

A motivating example

Query

“What symbolic representation does Dash et al. use for molecules in their neurosymbolic drug discovery framework?”

LLM-only response (typical):

“Dash et al. likely represent molecules using SMILES strings fed into a graph neural network. . .”

- ▶ Sounds reasonable but may be entirely fabricated
 - ▶ No source, no verifiability
 - ▶ Worse: the paper may postdate the model’s training cutoff
- ⇒ We need a way to supply the LLM with *relevant documents at query time.*

The RAG idea

Key insight: instead of compressing all knowledge into θ , maintain an *external knowledge base \mathcal{K}* and retrieve from it on demand.

Standard LLM

$$y^* = \arg \max_y P(y \mid q; \theta)$$

Knowledge source: θ only

RAG

$$y^* = \arg \max_y P(y \mid q, \mathcal{D}_q; \theta)$$

Knowledge sources: θ and \mathcal{D}_q

where $\mathcal{D}_q = \text{RETRIEVE}(q, \mathcal{K}, k)$ is a set of k document chunks retrieved from \mathcal{K} .

Lewis et al., *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*, NeurIPS 2020

Two phases of RAG

Phase 1: indexing (offline)

Build a searchable index over the knowledge base \mathcal{K} once.

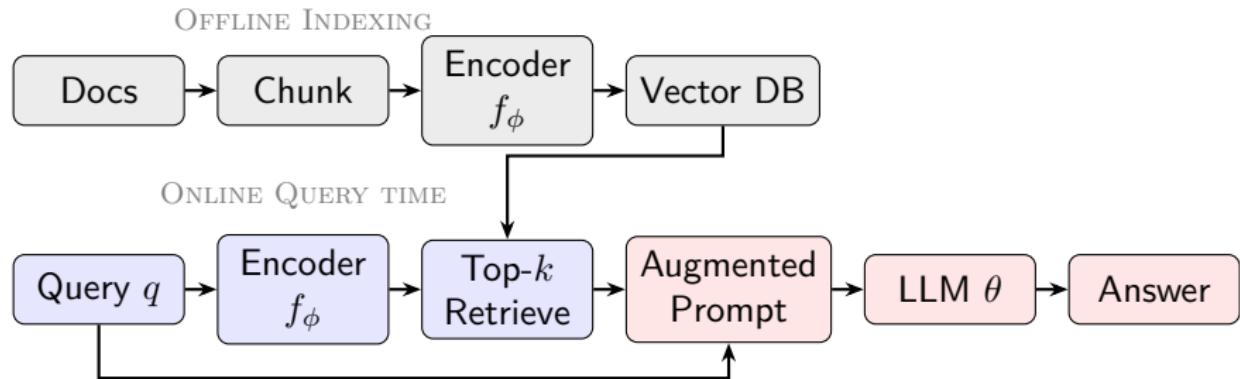
- (i) **Chunking**: split documents into manageable passages $\{d_1, d_2, \dots, d_N\}$
- (ii) **Encoding**: map each chunk to a dense vector $\mathbf{e}_i = f_\phi(d_i) \in \mathbb{R}^m$
- (iii) **Storing**: persist $\{\mathbf{e}_i\}$ in a vector database

Phase 2: query (online)

At inference time, for each user query q :

- (i) **Retrieve**: find top- k chunks most similar to q
- (ii) **Augment**: prepend retrieved chunks to the prompt
- (iii) **Generate**: pass augmented prompt to the LLM

RAG pipeline: overview



Augmented prompt:

$$\underbrace{[\text{System instruction}]}_{\text{role/task}} + \underbrace{[\mathcal{D}_q]}_{\text{retrieved chunks}} + \underbrace{[q]}_{\text{user query}}$$

Dense retrieval: semantic embeddings

The encoder f_ϕ maps text to a dense vector in \mathbb{R}^m such that *semantically similar* texts are geometrically close.

Cosine similarity between query q and chunk d_i :

$$\text{sim}(q, d_i) = \frac{f_\phi(q)^\top f_\phi(d_i)}{\|f_\phi(q)\| \|f_\phi(d_i)\|} \in [-1, 1]$$

Retrieval: return the k chunks with highest similarity

$$\mathcal{D}_q = \arg \max_{S \subseteq \mathcal{K}, |S|=k} \sum_{d_i \in S} \text{sim}(q, d_i)$$

In practice: **FAISS** (Facebook AI Similarity Search) enables approximate nearest-neighbour search at scale.

Johnson et al., IEEE TBDM 2019

Sparse vs. dense vs. hybrid retrieval

Property	Sparse (BM25)	Dense	Hybrid
Representation	TF-IDF weights	\mathbb{R}^m vector	Both
Matching	Exact term	Semantic	Both
Paraphrase recall	Poor	Good	Good
Rare terms	Good	Can miss	Good
Compute	Fast	Moderate	Higher

Reciprocal Rank Fusion (RRF) for hybrid scoring:

$$\text{score}_{\text{RRF}}(d_i) = \frac{1}{r_{\text{sparse}}(d_i) + c} + \frac{1}{r_{\text{dense}}(d_i) + c}$$

where $r(\cdot)$ is rank and c is a smoothing constant (typically 60).

Chunking strategy matters

The granularity of $\{d_1, \dots, d_N\}$ directly affects retrieval quality.

Fixed-size	Sentence-based	Recursive / hierarchical
Split every w words with overlap δ . Simple but may cut sentences.	Chunk at sentence boundaries. Semantically cleaner; variable length.	Respects document structure (sections, paragraphs). Preserves context best.

Chunk size trade-off:

- ▶ *Too small*: retrieved chunk lacks context \Rightarrow poor generation
- ▶ *Too large*: floods the LLM context window; retrieval precision drops

Conditional generation

Given retrieved chunks $\mathcal{D}_q = \{d_1^*, \dots, d_k^*\}$, the LLM generates:

$$y^* = \arg \max_y P(y \mid q, d_1^*, \dots, d_k^*; \theta)$$

Using the chain rule, this factorises token-by-token as before:

$$P(y \mid q, \mathcal{D}_q; \theta) = \prod_{t=1}^{|y|} P(y_t \mid y_{<t}, q, \mathcal{D}_q; \theta)$$

Nothing changes in the LLM itself.

The retrieved chunks are simply prepended to the input; θ remains fixed during inference.

⇒ RAG is a *prompting strategy*, not a fine-tuning strategy.

The context window constraint

The LLM can only process a finite number of tokens at once (the *context window*).

Budget decomposition:

$$L_{\text{total}} = L_{\text{system}} + \underbrace{\sum_{i=1}^k |d_i^*|}_{\text{retrieved context}} + L_q + L_y \leq L_{\max}$$

Increasing k gives more grounding but:

- ▶ Consumes more of the context window
- ▶ Risks burying the relevant chunk among noise (*lost-in-the-middle* problem)
- ▶ Increases latency

Practical choice: $k \in \{3, 5\}$ for most research QA tasks.

Liu et al., *Lost in the Middle: How LMs Use Long Contexts*, TACL 2024

Reranking: improving retrieval precision

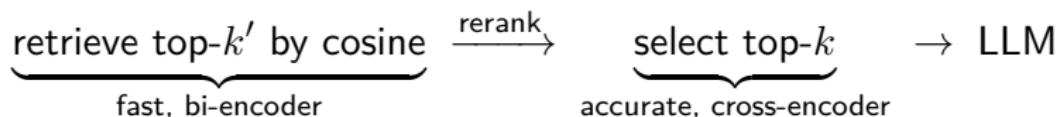
First-stage retrieval (cosine similarity) is fast but imprecise.

Reranker: a cross-encoder that scores each (q, d_i) pair jointly:

$$s_i = g_\psi(q, d_i) \in \mathbb{R}$$

- ▶ Sees *both* query and passage together \Rightarrow much more accurate
- ▶ Too slow to run over all N chunks; applied only to top- k' from first stage

Two-stage pipeline:



Nogueira & Cho, *Passage Re-ranking with BERT*, arXiv 2019

RAG vs. fine-tuning vs. long context

Property	RAG	Fine-tune	Long context
Updates knowledge	Dynamic	Static	Dynamic
Trains θ	No	Yes	No
Scales to large K	Yes	N/A	No
Attribution	Yes	No	Partial
Cost (inference)	Low	Low	High
Hallucination	Reduced	Moderate	Moderate

RAG is the preferred approach when the knowledge base is *large*, *dynamic*, or requires *traceable sources*.

How do we evaluate RAG?

Two sub-systems to evaluate independently and end-to-end:

Retrieval quality:

- ▶ *Recall@k*: fraction of relevant chunks in top- k retrieved
- ▶ *Mean Reciprocal Rank (MRR)*:

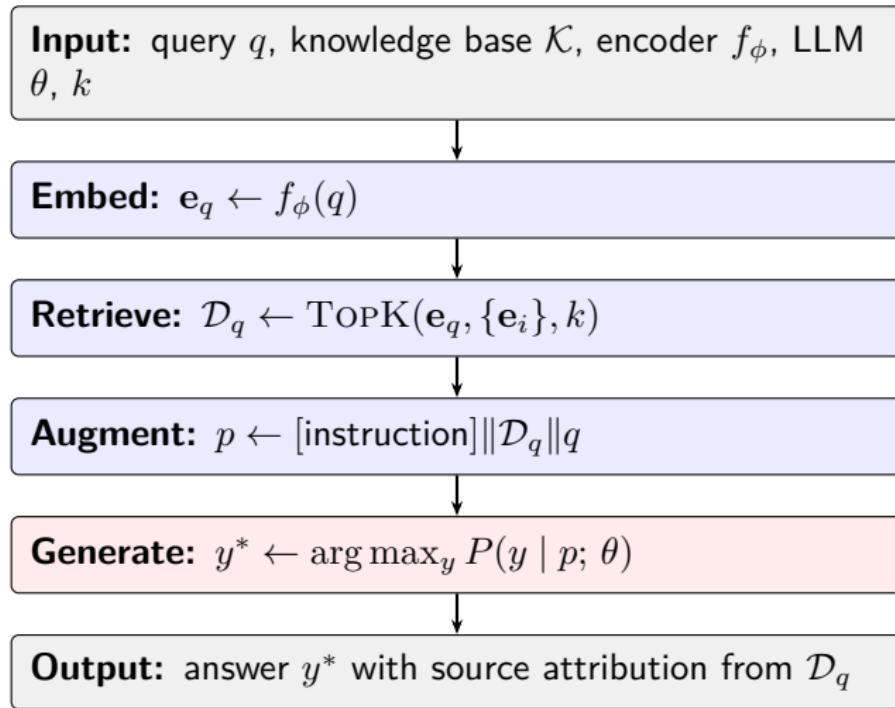
$$\text{MRR} = \frac{1}{|Q|} \sum_q \frac{1}{\text{rank of first relevant chunk}}$$

Generation quality (RAGAS framework):

- ▶ *Faithfulness*: is the answer supported by \mathcal{D}_q ?
- ▶ *Answer relevancy*: does the answer address q ?
- ▶ *Context recall*: do retrieved chunks cover the ground-truth answer?

Es et al., *RAGAS: Automated Evaluation of RAG*, EACL 2024
(arXiv:2309.15217)

RAG inference



Demo: NeSy QA (click [here](#))

Knowledge base \mathcal{K} : published papers on neurosymbolic AI and drug discovery.

Queries of interest:

- ▶ “*What symbolic representation is used for molecules?*”
- ▶ “*Which datasets are used for drug-likeness prediction?*”
- ▶ “*How is background knowledge encoded in the ILP framework?*”
- ▶ “*What is the performance vs. GNN baselines?*”

Comparison: **LLM-only** answer **vs.** **RAG-augmented** answer with source chunk displayed.

We will observe: LLM hallucinates method details; RAG cites exact text from the correct paper.